



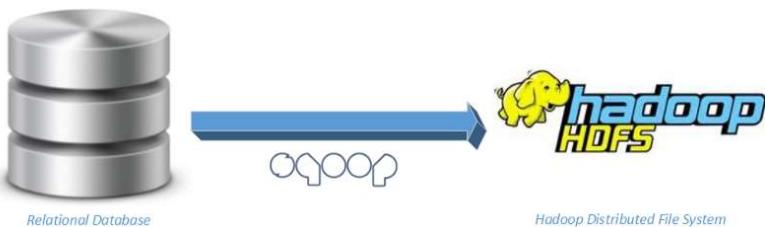
Tutorial 3: Apache Sqoop with Cloudera

CN7031 - Big Data Analytics

Dr Fahimeh Jafari (f.jafari@uel.ac.uk)

LEARNING OUTCOMES: After completing this tutorial, you should:

- Have gotten a hands-on experience in deploying Apache Sqoop with Cloudera
- Understand Sqoop concepts and commands
- Be able to import data into HDFS
- Using Hue to Verify/Query Tables



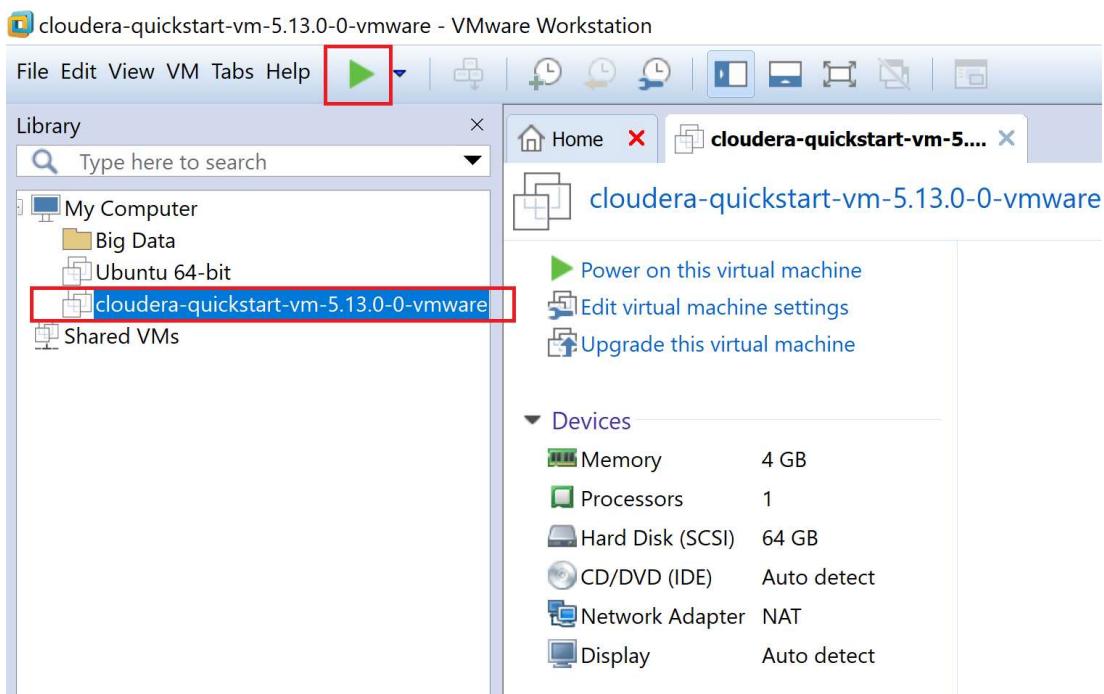
NOTE: Tutorial Submission (OPTIONAL)

You can submit the results of your work by taking screenshots (wherever pointed), pasting them into a word file and sending the word/pdf file through the submission link provided in Moodle.

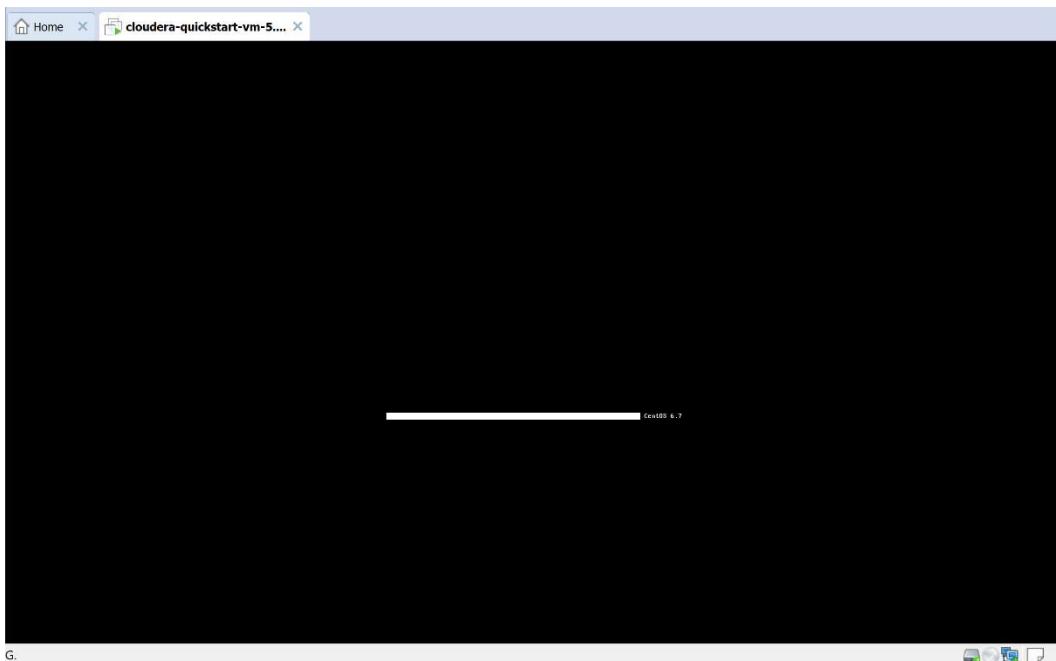
Exercise 1: Launch Cloudera

Please use the following instructions to launch the Cloudera Quickstart VM with VMware.

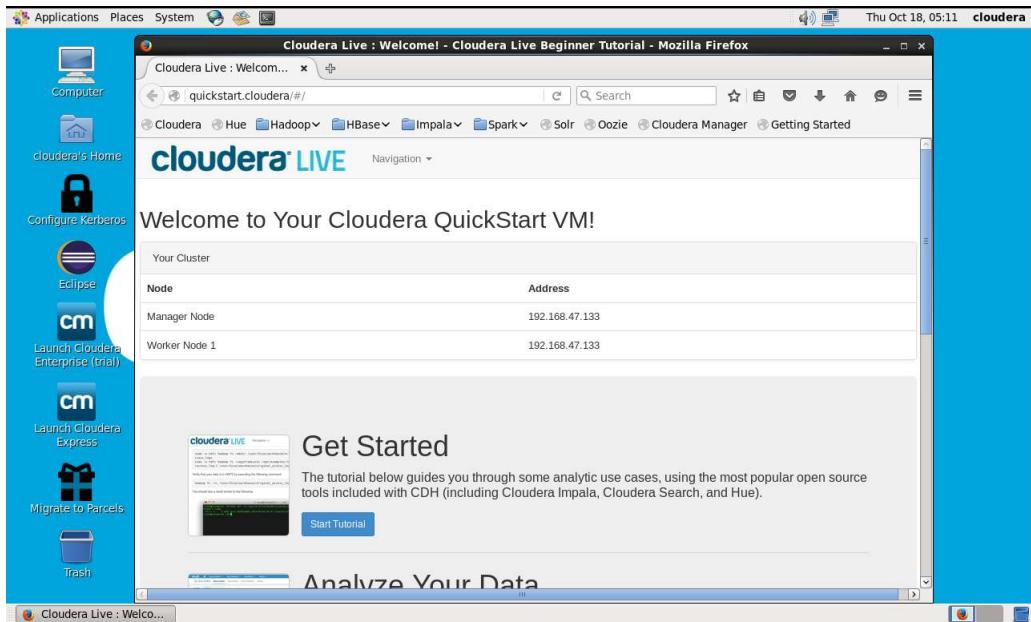
1. **Launch Cloudera VM.** When you open it, the `cloudera-quickstart-vm-5.13.0-0-vmware` will appear on the left in the VMware window. Select it and click the Start button to launch the VM.



2. **Cloudera VM booting.** It will take several minutes for the Virtual Machine to start. The booting process takes a long time since many Hadoop tools are started.



3. **The Cloudera VM desktop.** Once the booting process is complete, the desktop will appear with a browser.



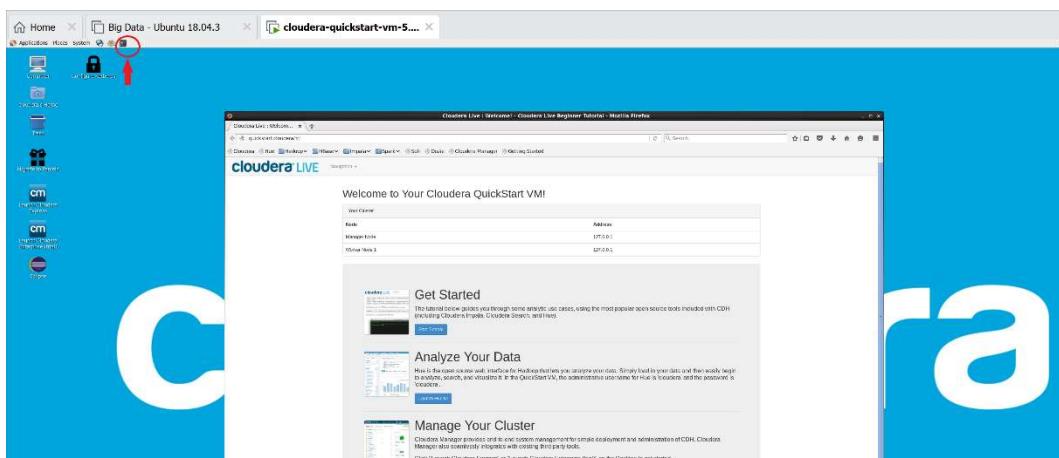
Exercise 2: Work with MySQL

MySQL is an open-source database management system, commonly installed as part of the popular LAMP (Linux, Apache, MySQL, PHP/Python/Perl) stack. It uses a relational database and SQL (Structured Query Language) to manage its data.

We don't need to install MySQL package as it's installed in Cloudera image.

Step 1: How to Access the MySQL shell

1. You should first open a terminal, which you can do by clicking on the black "Terminal" icon at the top of your screen. Once it is open, you can launch the MySQL.



2. Run MySQL by typing the following command:



```
1 mysql -u root -p
```

3. Then you'll be asked to enter MySQL password. Type **cloudera**

You'll then be presented with the MySQL monitor prompt:

A screenshot of a terminal window titled "cloudera@quickstart:~". The window shows the MySQL monitor interface. The text output includes the MySQL prompt "[cloudera@quickstart ~]\$ mysql -u root -p", the password entry prompt "Enter password:", the welcome message "Welcome to the MySQL monitor. Commands end with ; or \g.", the connection ID "Your MySQL connection id is 298", the server version "Server version: 5.1.73 Source distribution", copyright information "Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.", a trademark notice "Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.", and help instructions "Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.". The MySQL prompt "mysql>" is visible at the bottom.

Two points to keep in mind:

- All MySQL commands end with a semicolon; if the phrase does not end with a semicolon, the command will not execute.
- Also, although it is **not** required, MySQL commands are usually written in uppercase and databases, tables, usernames, or text are in lowercase to make them easier to distinguish. However, the MySQL command line is not case sensitive.

Step 2: Show Database and Tables

MySQL organizes its information into databases; each one can hold tables with specific data. Show existing databases with

```
SHOW DATABASES;
```

Output



```
mysql> SHOW DATABASES;
+-----+
| Database |
+-----+
| information_schema |
| cm |
| firehose |
| hue |
| metastore |
| mysql |
| nav |
| navms |
| oozie |
| retail_db |
| rman |
| sentry |
+-----+
12 rows in set (0.00 sec)
```

(Take Screenshot #1)

In this tutorial we works with `retail_db` database So, access to this databases with

```
USE retail_db;
```

You can quickly check what tables are available by typing:

```
SHOW TABLES;
```

Output

```
mysql> SHOW TABLES;
+-----+
| Tables_in_retail_db |
+-----+
| categories |
| customers |
| departments |
| order_items |
| orders |
| products |
+-----+
6 rows in set (0.00 sec)
```

(Take Screenshot #2)

Step 3: Exit from MySQL

Type following command to exit from MySQL

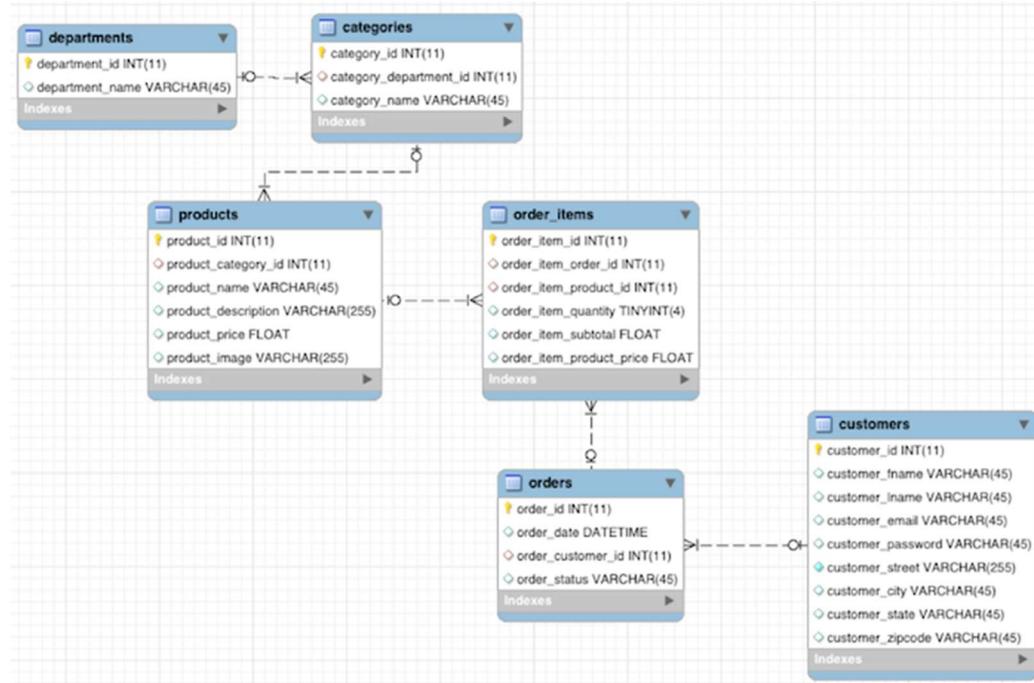
```
exit
```

Exercise 3: Importing data from MySQL database to Hadoop HDFS

To analyse the transaction data in the new platform, we need to ingest it into the Hadoop Distributed File System (HDFS). We need to find a tool that easily transfers structured data from a RDBMS to HDFS, while preserving structure. That enables us to query the data, but not interfere with or break any regular workload on it.



This exercise describes how to import data from MySQL database to Hadoop HDFS. We are going to import `retail_db` into the Hadoop with Sqoop. The flowchart of `retail_db` is in the following picture [Source: <https://www.cloudera.com/developers/get-started-with-hadoop-tutorial/exercise-1.html>]



Open a new terminal and run the following Sqoop command line by line. [It would take a while to be completed, because Sqoop launches MapReduce job.](#)

NOTE: You must type the commands yourself. Don't copy & past them as it wouldn't work.

```
[cloudera@quickstart ~]$ sqoop import-all-tables \
-m 2 \
--connect jdbc:mysql://localhost:3306/retail_db \
--username=root \
--password=cloudera \
--compression-codec=snappy \
--as-parquetfile \
--warehouse-dir /user/hive/warehouse
```

After about 1 minute, you have to see the following screenshot as the confirmation:



The screenshot shows a terminal window titled "cloudera@quickstart:~". The output displays various HDFS and MapReduce counters. It includes sections for Job Counters, Map-Reduce Framework, File Input Format Counters, and File Output Format Counters. Log entries from mapreduce.ImportJobBase are also present, indicating file transfers and record retrievals. The terminal prompt "[cloudera@quickstart ~]\$ " is at the bottom.

```
HDFS: Number of write operations=20
Job Counters
    Launched map tasks=2
    Other local map tasks=2
    Total time spent by all maps in occupied slots (ms)=27937
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=27937
    Total vcore-milliseconds taken by all map tasks=27937
    Total megabyte-milliseconds taken by all map tasks=28607488
Map-Reduce Framework
    Map input records=1345
    Map output records=1345
    Input split bytes=236
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=532
    CPU time spent (ms)=9040
    Physical memory (bytes) snapshot=857354240
    Virtual memory (bytes) snapshot=3157467136
    Total committed heap usage (bytes)=711983104
File Input Format Counters
    Bytes Read=0
File Output Format Counters
    Bytes Written=0
19/10/20 07:34:12 INFO mapreduce.ImportJobBase: Transferred 55.624 KB in 112.2176 seconds (507.5763 bytes/sec)
19/10/20 07:34:12 INFO mapreduce.ImportJobBase: Retrieved 1345 records.
[cloudera@quickstart ~]$
```

(Take Screenshot #3)

Verification:

When this command is complete, confirm that your data files exist in HDFS. These commands will show the directories and the files inside them that make up your tables:

```
hdfs dfs -ls /user/hive/warehouse/
hdfs dfs -ls /user/hive/warehouse/categories/
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /user/hive/warehouse/
Found 6 items
drwxrwxrwx  - cloudera supergroup      0 2019-10-20 14:57 /user/hive/warehouse/categories
drwxrwxrwx  - cloudera supergroup      0 2019-10-20 14:59 /user/hive/warehouse/customers
drwxrwxrwx  - cloudera supergroup      0 2019-10-20 15:02 /user/hive/warehouse/departments
drwxrwxrwx  - cloudera supergroup      0 2019-10-20 15:04 /user/hive/warehouse/order_items
drwxrwxrwx  - cloudera supergroup      0 2019-10-20 15:06 /user/hive/warehouse/orders
drwxrwxrwx  - cloudera supergroup      0 2019-10-20 15:08 /user/hive/warehouse/products
[cloudera@quickstart ~]$ hdfs dfs -ls /user/hive/warehouse/categories
Found 4 items
drwxr-xr-x  - cloudera supergroup      0 2019-10-20 14:56 /user/hive/warehouse/categories/.metadata
drwxr-xr-x  - cloudera supergroup      0 2019-10-20 14:57 /user/hive/warehouse/categories/.signals
-rw-r--r--  1 cloudera supergroup 1491 2019-10-20 14:57 /user/hive/warehouse/categories/0e01f6e7-4f7a-4b66-9225-34c2bbe9bfd6.parquet
-rw-r--r--  1 cloudera supergroup 1520 2019-10-20 14:57 /user/hive/warehouse/categories/c141345a-2a78-48cc-8801-058e7e53d102.parquet
```

(Take Screenshot #4)

Note: The number of .parquet files shown will be equal to what was passed to Sqoop with the -m parameter. This is the number of 'mappers' that Sqoop will use in its MapReduce jobs. It could also be thought of as the number of simultaneous connections to your database, or the number of disks/Data Nodes you want to spread the data across. So, on a single-node you will just see one, but larger clusters will have a greater number of files.

OR, complete the verification step visually (localhost:50070) through the following steps:

- Open a new tab in Firefox and go to <http://localhost:50070/>



- Then click on Utilities->Brows the file systems

Namenode Information - Mozilla Firefox
Cloudera Live : Welcome... Namenode information
localhost:50070/dshealth.html#tab-overview
Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started
Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities
Logs Browse the file system

Overview 'quickstart.cloudera:8020' (active)

Started:	Mon Oct 29 06:42:08 -0700 2018
Version:	2.6.0-cdh5.13.0, r42e8860b182e55321bd5f5605264da4adc8882be
Compiled:	Wed Oct 04 11:08:00 -0700 2017 by jenkins from Unknown
Cluster ID:	CID-768feab0-9bd0-4e04-a0ae-8787b64d9475
Block Pool ID:	BP-333635372-127.0.0.1-1508779710286

Summary

Security is off.
Safemode is off.
1,197 files and directories, 1,025 blocks = 2,222 total filesystem object(s).
Heap Memory used 65.64 MB of 350 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 36.54 MB of 37.44 MB Committed Non Heap Memory. Max Non Heap Memory is 130 MB.

Configured Capacity:	54.51 GB
DFS Used:	844.9 MB (1.51%)
Non DFS Used:	8.1 GB
DFS Remaining:	42.56 GB (78.08%)
Block Pool Used:	844.9 MB (1.51%)

You should see user Directory.

Browsing HDFS - Mozilla Firefox
Cloudera Live : Welcome... Browsing HDFS
localhost:50070/explorer.html#/ Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started
Hadoop Overview Datanodes Snapshot Startup Progress Utilities

Browse Directory

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxrwxrwx	hdfs	supergroup	0 B	Mon Oct 23 10:29:49 -0700 2017	0	0 B	benchmarks
drwxr-xr-x	hbase	supergroup	0 B	Fri Oct 16 03:06:13 -0700 2020	0	0 B	hbase
drwxr-xr-x	solr	solr	0 B	Mon Oct 23 10:32:05 -0700 2017	0	0 B	solr
drwxrwxrwt	hdfs	supergroup	0 B	Fri Oct 16 03:27:17 -0700 2020	0	0 B	tmp
drwxr-xr-x	hdfs	supergroup	0 B	Mon Oct 23 10:31:37 -0700 2017	0	0 B	user
drwxr-xr-x	hdfs	supergroup	0 B	Mon Oct 23 10:31:29 -0700 2017	0	0 B	var

Hadoop, 2017.

- Click on it and then go to hive-> warehouse. Now you should see the directories and the files that make up your tables:

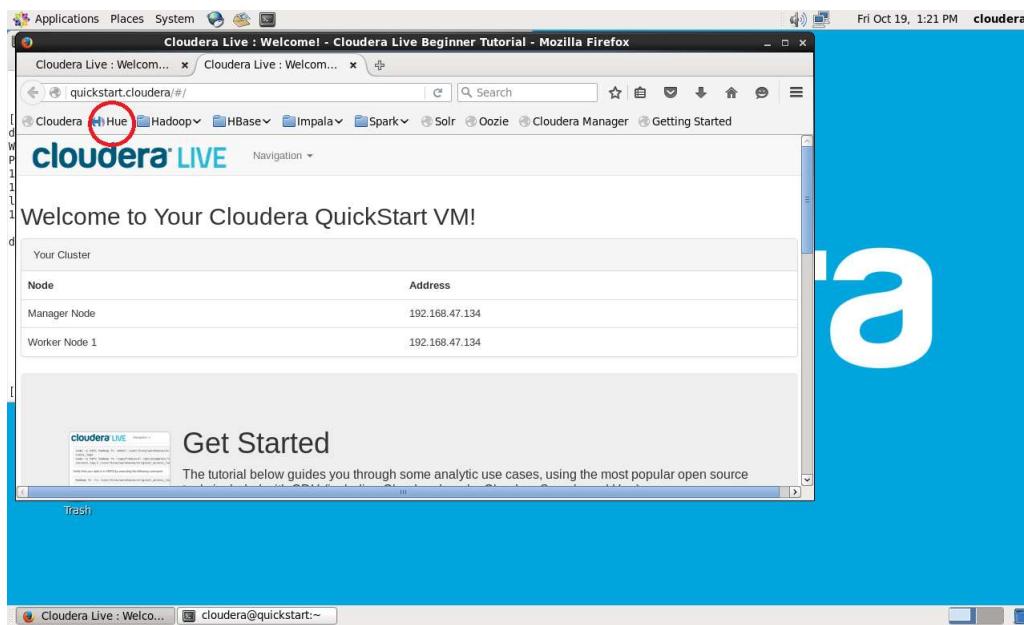


/user/hive/warehouse								Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
drwxrwxrwx	cloudera	supergroup	0 B	Sun Oct 20 14:57:39 -0700 2019	0	0 B	categories	
drwxrwxrwx	cloudera	supergroup	0 B	Sun Oct 20 14:59:25 -0700 2019	0	0 B	customers	
drwxrwxrwx	cloudera	supergroup	0 B	Sun Oct 20 15:02:03 -0700 2019	0	0 B	departments	
drwxrwxrwx	cloudera	supergroup	0 B	Sun Oct 20 15:04:20 -0700 2019	0	0 B	order_items	
drwxrwxrwx	cloudera	supergroup	0 B	Sun Oct 20 15:06:39 -0700 2019	0	0 B	orders	
drwxrwxrwx	cloudera	supergroup	0 B	Sun Oct 20 15:08:46 -0700 2019	0	0 B	products	

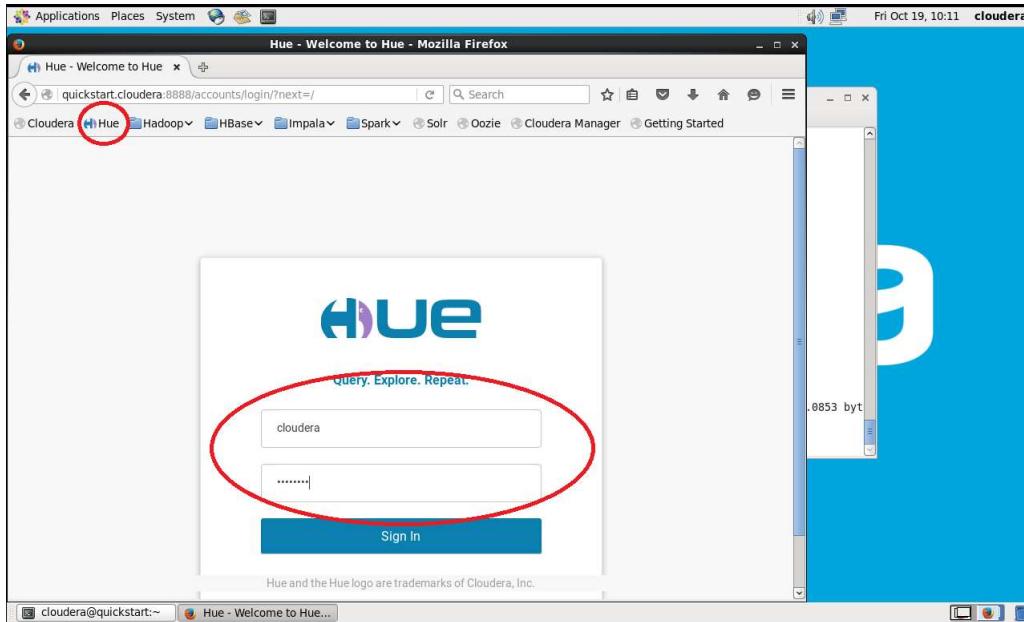
(Take Screenshot #5)

Exercise 4: Using Hue to Verify/Query Tables

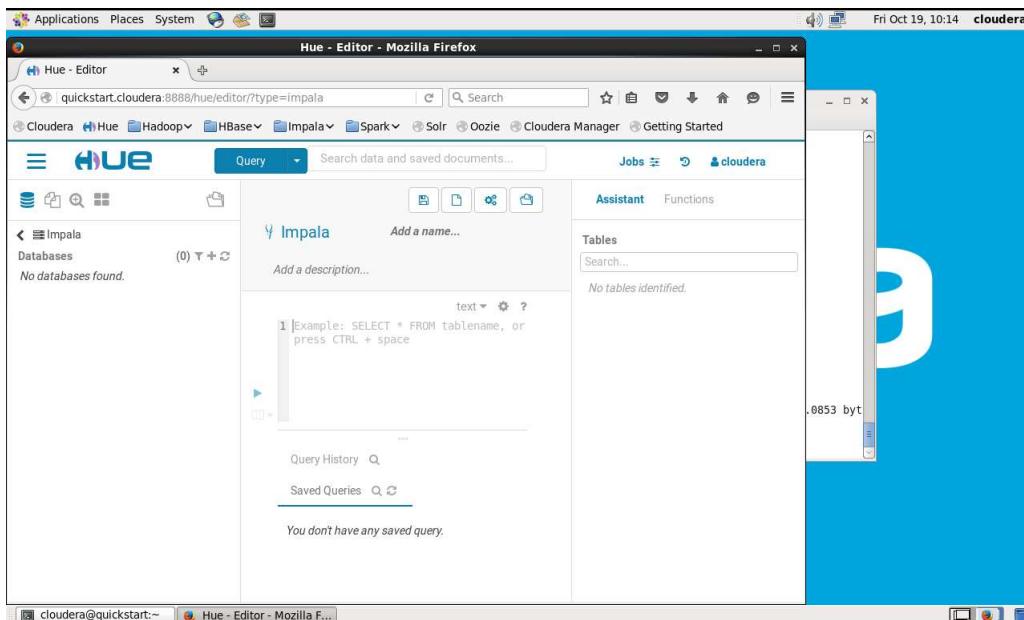
We're going to use Hue's Impala app to query our tables. Hue provides a web-based interface for many of the tools in CDH and can be found on port 8888 of your Manager Node. In the QuickStart VM, click on Hue as shown in the picture.



Then enter the administrator username for Hue which is 'cloudera'. For the password, please first try 'BigDataadmin', if it doesn't work, try 'cloudera'.



Once you are inside of Hue, click on **Query → Editor → Impala** to launch the Impala Editor.



To save time during queries, Impala does not poll constantly for metadata changes. So, the first thing we must do is tell Impala that its metadata is out of date. Then we should see our tables show up, ready to be queried. To this, type the following commands in the editor and then click on the play icon to run the commands.

```
invalidate metadata;  
  
show tables;
```



Impala

Add a name... Add a description...

1 show tables

Query History Saved Queries Results (6)

name
1 categories
2 customers
3 departments
4 order_items
5 orders
6 products

[Exceptional Error:] You may get the following error when you run the commands. This is because both services `impala-state-store` and `impala-catalog` are down and need to be restarted.

Hue - Editor - Mozilla Firefox

Fri Oct 19, 10:19 cloudera

Hue - Editor quickstart.cloudera:8888/hue/editor?editor=3

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

HUE Query Search data and saved documents...

Impala Add a name... Add a description... 1 INVALIDATE METADATA; 2 SHOW TABLES;

AnalysisException: This Impala daemon is not ready to accept user requests. Status: Waiting for catalog update from the StateStore.

Tables Search... No tables identified.

Jobs Assistant Functions

Query History Saved Queries

To fix the error, back to the terminal and type the following commands:

- `sudo service impala-state-store start`
- `sudo service impala-catalog start`



The screenshot shows a terminal window within the Hue interface. The terminal output includes commands like 'mapreduce.ImportJobBase', 'hadoop fs -ls /user/hive/warehouse/emp', and 'sudo service impala-state-store start'. A red box highlights the command 'sudo service impala-catalog start' and its response '[OK]'. Below the terminal, a message box displays: 'AnalysisException: This Impala daemon is not ready to accept user requests. Status: Waiting for catalog update from the StateStore.' At the bottom of the terminal window, there are 'Query History' and 'Saved Queries' links.

Now click on Hue icon again to refresh the page and then run the query. You can now see the tables in the result menu.

Now, your transaction data is readily available for structured queries in CDH, it's time to address a few questions. You can make a few basic/advanced SQL queries to get familiar with the `retail_db` database. Also, you can see the tables' properties visually in the Hue platform as shown in the following picture.

The screenshot shows the results of a 'show tables' query in the Hue interface. On the left, a sidebar lists tables: categories, customers, departments, order_items, orders, and products. The main area shows the query '1 show tables' and its results. The results table has columns 'name' and 'id'. The data is:

name	id
categories	1
customers	2
departments	3
order_items	4
orders	5
products	6

Type and run the following SQL query in the editor to show the content of the `customer` table:

```
SELECT * FROM customers;
```

It shows you the `customer` table data in the result window.



The screenshot shows the Hue Editor interface. On the left, there's a sidebar with a tree view of tables: default, categories, customers, departments, order_items, orders, and products. The 'customers' table is selected. In the main area, a query is run: 'SELECT * FROM customers;'. The results show 1,024 rows with columns: customer_id, customer_fname, customer_lname, customer_email, customer_password, customer_street, customer_city, and customer_state. The data includes various names like Richard Hernandez, Mary Barrett, Ann Smith, etc., along with their addresses and city-state information.

customer_id	customer_fname	customer_lname	customer_email	customer_password	customer_street	customer_city	customer_state
1	Richard	Hernandez	XXXXXXXXXX	XXXXXXXXXX	6303 Heather Plaza	Brownsville	TX
2	Mary	Barrett	XXXXXXXXXX	XXXXXXXXXX	9526 Noble Embers Ridge	Littleton	CO
3	Ann	Smith	XXXXXXXXXX	XXXXXXXXXX	3422 Blue Pioneer Bend	Caguas	PR
4	Mary	Jones	XXXXXXXXXX	XXXXXXXXXX	8324 Little Common	San Marcos	CA
5	Robert	Hudson	XXXXXXXXXX	XXXXXXXXXX	10 Crystal River Mall	Caguas	PR
6	Mary	Smith	XXXXXXXXXX	XXXXXXXXXX	3151 Sleepy Quail Promenade	Passaic	NJ
7	Melissa	Wilcox	XXXXXXXXXX	XXXXXXXXXX	9453 High Concession	Caguas	PR
8	Megan	Smith	XXXXXXXXXX	XXXXXXXXXX	3047 Foggy Forest Plaza	Lawrence	MA
9	Mary	Perez	XXXXXXXXXX	XXXXXXXXXX	3616 Quaking Street	Caguas	PR
10	Melissa	Smith	XXXXXXXXXX	XXXXXXXXXX	8598 Harvest Beacon Plaza	Stafford	VA
11	Mary	Huffman	XXXXXXXXXX	XXXXXXXXXX	3169 Stony Woods	Caguas	PR
12	Christopher	Smith	XXXXXXXXXX	XXXXXXXXXX	5594 Jagged Embers By-pass	San Antonio	TX
13	Mary	Baldwin	XXXXXXXXXX	XXXXXXXXXX	7922 Iron Oak Gardens	Caguas	PR
14	Katherine	Smith	XXXXXXXXXX	XXXXXXXXXX	5666 Hazy Pony Square	Pico Rivera	CA
15	Jane	Luna	XXXXXXXXXX	XXXXXXXXXX	673 Burning Glen	Fontana	CA
16	Tiffany	Smith	XXXXXXXXXX	XXXXXXXXXX	6651 Iron Port	Caguas	PR
17	Mary	Robinson	XXXXXXXXXX	XXXXXXXXXX	1325 Noble Pike	Taylor	MI
18	Robert	Smith	XXXXXXXXXX	XXXXXXXXXX	2734 Hazy Butterfly Circle	Martinez	CA
19	Stephanie	Mitchell	XXXXXXXXXX	XXXXXXXXXX	3543 Red Treasure Bay	Caguas	PR
20	Mary	Ellis	XXXXXXXXXX	XXXXXXXXXX	4703 Old Route	West New York	NJ
21	William	Johnson	XXXXXXXXXX	XXXXXXXXXX	9999 Main Street	Phoenix	AZ

(Take Screenshot #6)

Extra Task: Exporting data from Hadoop to MySQL

You can use the export command to export data back from the HDFS to MySQL.

```
$ sqoop export --generic-args --export-args  
$ sqoop-export --generic-args --export-args
```

The target table must exist in the target database. The files that are given as input to the Sqoop contain records, which are called rows in the table. Those are read and parsed into a set of records and delimited with a user-specified delimiter. The default operation is to insert all the records from the input files into the database table using the INSERT statement. In update mode, Sqoop generates the UPDATE statement that replaces the existing record in the database. It only understands HDFS directories, not Hive or Hcatalog.

NOTE: It is mandatory that the table to be exported is created manually and is present in the MySQL database.

Let us take the customer's data in the file, in HDFS. The customer's data is available in the retail file in the retail_db directory in HDFS. It is mandatory that the table to be exported is created manually and is present in the database from where it has to be exported. So, now we are creating a table named customers_export in MySQL and SELECT all the data from the customers table.



NOTE: You must type the commands yourself. Don't copy and paste them as it wouldn't work.

1. Either keep the mysql terminal open, or issue a new terminal and invoke the following;

```
mysql> use retail_db;
```

2. Create a new table “customers_export” to start the importing / exporting from MySQL to HDFS

```
CREATE TABLE customers_export AS  
  
SELECT  
  
    customer_id,  
  
    customer_fname,  
  
    customer_lname,  
  
    customer_email,  
  
    customer_password,  
  
    customer_street,  
  
    customer_city,  
  
    customer_state,  
  
    customer_zipcode  
  
FROM customers  
  
WHERE 1=0;
```

3. Verify the new table “customers_export” exists;

```
mysql> show tables;
```

4. Verify it's currently empty;

```
mysql> select * from customers_export;
```



5. Now we import from MySQL to HDFS

```
[cloudera@quickstart ~]$ sqoop import \
--connect jdbc:mysql://quickstart.cloudera:3306/retail_db \
--username root \
--password cloudera \
--table customers \
--target-dir /retail_db/customers
```

After about 3 minutes, you must see the following screenshot as the confirmation:

The screenshot shows a VMware Workstation Player window titled "cloudera-quickstart-vm-5.13.0-0-vmware - VMware Workstation 17 Player (Non-commercial use only)". Inside the window, there is a terminal window titled "cloudera@quickstart:~" displaying Java stack trace output from a Sqoop job. To the right of the terminal is a "Hue Manager" interface showing a "Tables" section with a search bar and a message "No tables identified". Below the terminal and Hue interface is a table view showing customer data:

ID	Customer Name	Role
2	Walter	SI
3	Mary	SI
4	Andrew	M
5	Stephanie	SI
6	Carolyn	St
7	Scott	H
8	Mary	SI



(Take Screenshot #7)

6. Now we export from HDFS into MySQL to fill in the data of the new table “customers_export”

```
[cloudera@quickstart ~]$ sqoop export \
--table customers_export \
--connect jdbc:mysql://quickstart.cloudera:3306/retail_db \
--username root \
--password cloudera \
--export-dir /retail_db/customers
```

7. Verify export transaction by issuing the following in the same open “MySQL” terminal.

```
mysql> select * FROM customers_export;
```

Or, open a new mysql terminal to verify exportation, and then follow the following instruction to verify.

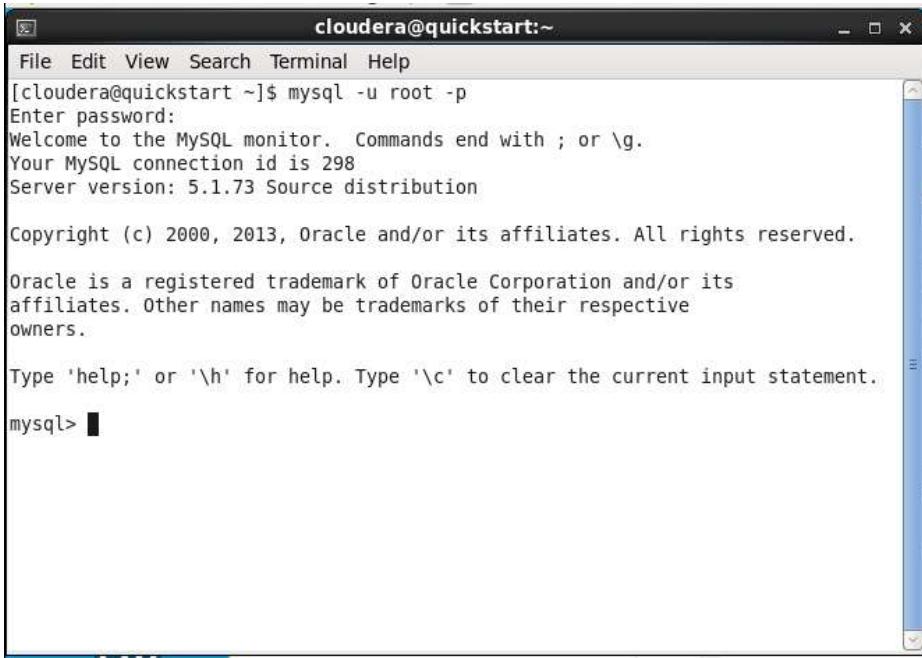
Run MySQL by typing the following command:

```
1 mysql -u root -p
```

Then you'll be asked to enter MySQL password. Type **cloudera**



You'll then be presented with the MySQL monitor prompt:



The screenshot shows a terminal window titled "cloudera@quickstart:~". The window displays the MySQL monitor interface. The text output includes:

```
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 298
Server version: 5.1.73 Source distribution

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> █
```

To get access to these databases invoke the following:

```
USE retail_db;
```



The following command is used to verify the table in the MySQL command line:

```
mysql> select * from customers_export;
```

After running this command, you will see the customer's data in MySQL

A screenshot of a terminal window titled "cloudera@quickstart:~". The window shows the output of a MySQL query. The output displays a table of customer data with columns: ID, First Name, Last Name, Street Address, City, State, and Zip Code. The zip codes are partially redacted with 'XXXXXX'. The terminal window is part of a desktop environment with a menu bar and taskbar visible.

ID	First Name	Last Name	Street Address	City	State	Zip Code
12425	Mary	Smith	1050 Grand Forest Towers	Caguas	PR	XXXXXXX
00725			12426 Jordan	Valdez		XXXXXXX
			5561 Quiet Loop		NY	XXXXXXX
11210			12427 Mary	Smith		XXXXXXX
			3662 Round Barn Gate		TX	XXXXXXX
75093			12428 Jeffrey	Travis		XXXXXXX
			1552 Burning Dale Highlands	Caguas	PR	XXXXXXX
00725			12429 Mary	Smith		XXXXXXX
			92 Sunny Bear Villas	Gardena	CA	XXXXXXX
90247			12430 Hannah	Brown		XXXXXXX
			8316 Pleasant Bend	Caguas	PR	XXXXXXX
00725			12431 Mary	Rios		XXXXXXX
			1221 Cinder Pines		HI	XXXXXXX
96744			12432 Angela	Smith		XXXXXXX
			1525 Jagged Barn Highlands	Caguas	PR	XXXXXXX
00725			12433 Benjamin	Garcia		XXXXXXX
			5459 Noble Brook Landing	Levittown	NY	XXXXXXX
11756			12434 Mary	Mills		XXXXXXX
			9720 Colonial Parade	Caguas	PR	XXXXXXX
00725			12435 Laura	Horton		XXXXXXX
			5736 Honey Downs		SC	XXXXXXX
29483						

```
+-----+-----+-----+-----+-----+
| 12435 rows in set (0.06 sec)
```

```
mysql> [cloudera@quickstart:~]
```

The terminal window has a title bar "cloudera@quickstart:~". The taskbar at the bottom shows several open windows, including "Hue - Editor - Mo...", "cloudera@quicks...", "cloudera@quicks...", "cloudera@quicks...", and "cloudera@quicks...".