



University of  
East London

Pioneering Futures Since 1898

# How Businesses Work with Unstructured Data Using Hadoop

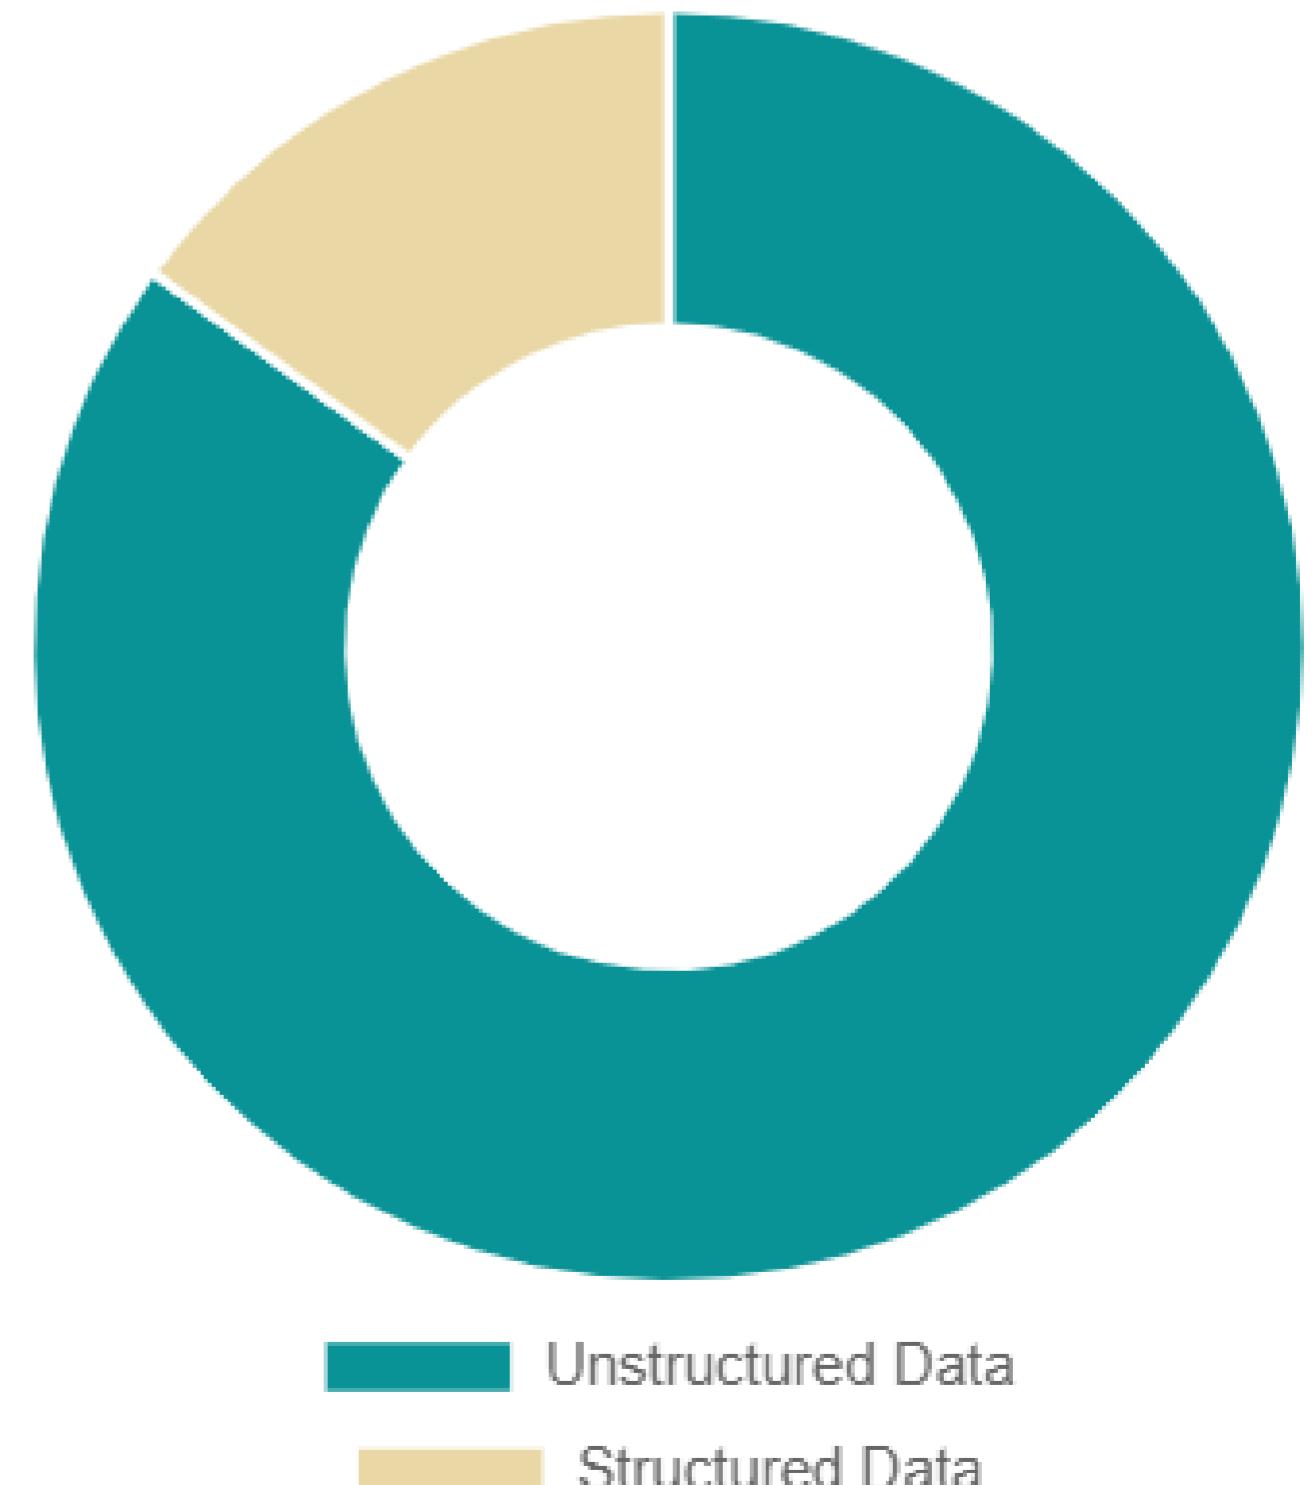
Dr Amin Karami

Associate Professor in AI

[a.karami@uel.ac.uk](mailto:a.karami@uel.ac.uk)

[www.aminkarami.com](http://www.aminkarami.com)

CN7031 – Lecture  
October 2025 – Week 5





## Speaker: Setinaz Foroudi

From Challenges to Opportunities, a Journey of Growth and Innovation

🌟 Gain insights from an **industry expert** on transitioning from academia to business

- 📅 Wednesday (19 Nov): MLT - 16:45-18:00
- 📅 Thursday (20 Nov): WB.G.02 - 16:45-18:00

Expertise: AI, Data Analytics, Learning Systems



University  
of  
East London

Pioneering Futures since 1898

# Outline

- The Hadoop Workflow for Unstructured Data
- Schema-on-Read and Schema-on-Write
- How Hive Reads Data: The "SerDe"
- Advanced Hive Optimizations



University of  
East London

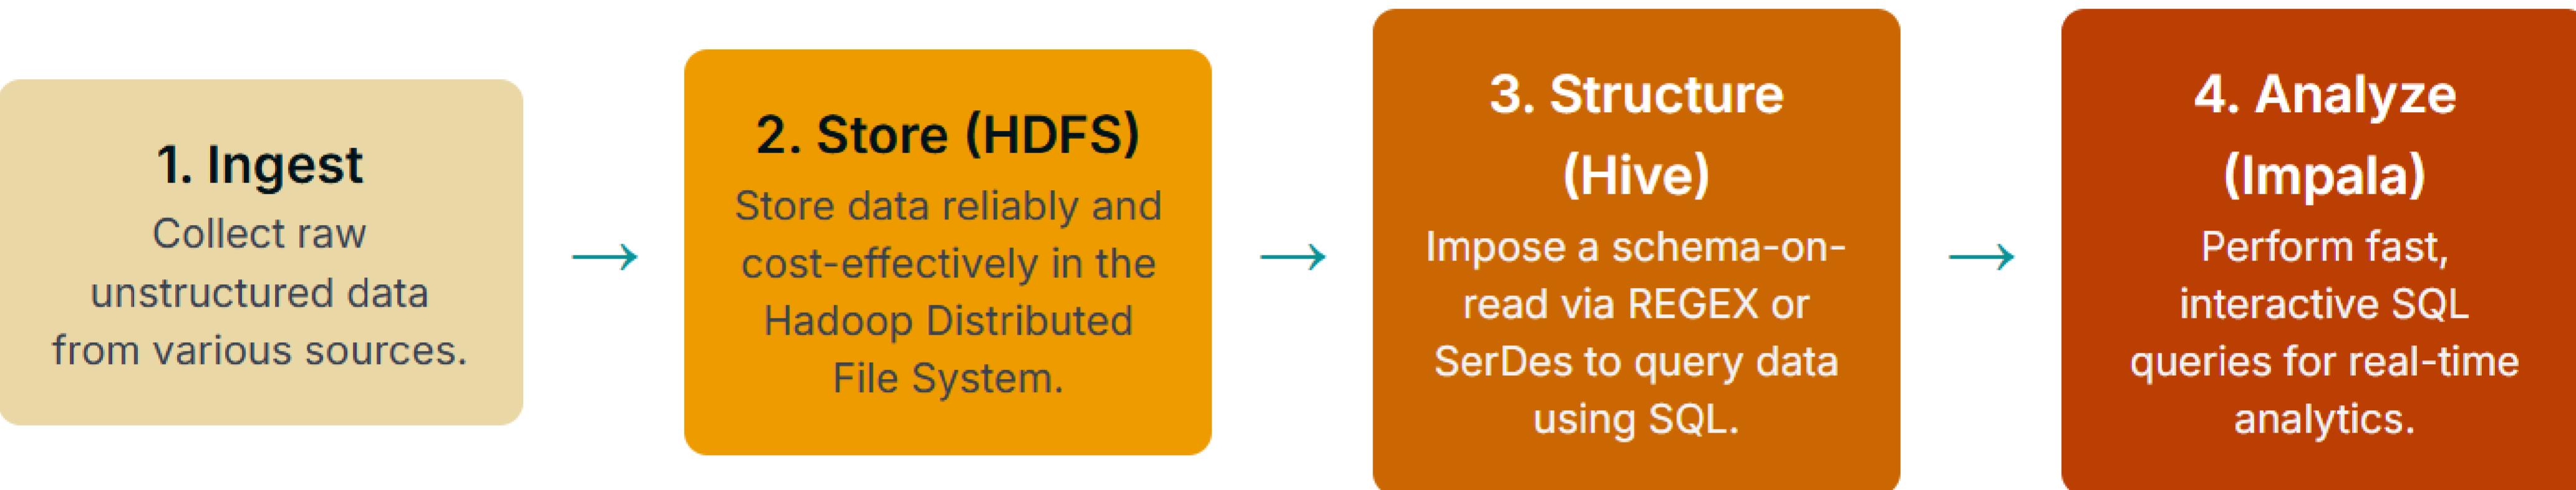
# Unstructured Data

Data that does not fit into a traditional row-column database. It is complex and comes in various formats:

- Social Media posts
- Log Files
- Images & Videos
- Sensor Data
- Emails



# The Hadoop Ecosystem Workflow



## Imposing Structure with Hive:

Hive's "schema-on-read" is the key to querying unstructured data with familiar SQL syntax.



University of  
East London

# Schema-on-read

- The schema is applied to the data only **at the time of query execution** (without upfront schema enforcement), allowing unstructured or semi-structured data, suitable for Big data systems and data lakes (e.g., Hive).

```
-- Create an external table for sales data
CREATE EXTERNAL TABLE sales_data (
    sale_id INT,          -- Unique ID for each sale
    sale_amount FLOAT,     -- Amount of the sale
    sale_date STRING       -- Date of the sale
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','  -- Data fields are separated by commas
STORED AS TEXTFILE
LOCATION '/data/sales/';  -- Location of the data files
```

```
-- Example rows in the data:
-- 1,100.50,2025-10-01
-- 2,200.75,2025-10-02

-- Query to retrieve sales with an amount greater than 150
SELECT sale_id, sale_amount
FROM sales_data
WHERE sale_amount > 150;
```



# Schema-on-write

- The schema is enforced **when data is written** to the storage system, ensuring that only data following the predefined schema is stored.
- This approach is used in traditional relational databases like MySQL or PostgreSQL, where data integrity and structure are critical.

```
CREATE TABLE sales_data (
    sale_id INT NOT NULL,
    sale_amount FLOAT NOT NULL,
    sale_date DATE NOT NULL
);
```

```
INSERT INTO sales_data (sale_id, sale_amount, sale_date)
VALUES (1, 100.50, '2025-10-01');

SELECT sale_id, sale_amount
FROM sales_data
WHERE sale_amount > 150;
```



# The Role of SerDe

SerDe is a core component of Hive's flexibility, enabling it to process data stored in a variety of formats. It supports the schema-on-read paradigm by dynamically interpreting data at query time.

## Popular SerDe Choices:

- **JSON SerDe**: parsing JSON files.
- **Regex SerDe**: applying regular expressions to unstructured text like log files.
- **Avro/Parquet/ORC SerDe**: efficient binary formats.
- **Multimedia SerDe**: multimedia files (audio, video, images)



# SerDe: Example

```
CREATE EXTERNAL TABLE intermediate_logs (
    ip STRING,
    datee STRING,
    method STRING,
    url STRING,
    http_version STRING,
    code1 STRING,
    code2 STRING,
    dash STRING,
    user_agent STRING
)
ROW FORMAT SERDE 'org.apache.hadoop.hive.contrib.serde2.RegexSerDe'
WITH SERDEPROPERTIES (
    'input.regex' = '([^\ ]*) - - \[(\d{2}/[A-Za-z]{3}/\d{4}:\d{2}:\d{2}:\d{2} -\d{4})...
    ... (complete regex pattern)...
)
LOCATION '/path/to/raw/logs';
```



# REGEX in action

## 1. Raw Unstructured Data:

```
127.0.0.1 - - [17/Oct/2025:17:15:00 +0100] "GET /index.html H  
TTP/1.1" 200 2326
```



## 2. Apply REGEX Pattern via SerDe:

```
^(\S+) (\S+) (\S+) \[(\w:/]+\s[+\-]\d{4})\] "(.+?)" (\d{3})  
(\d+)
```



## 3. Result: A Queryable "Virtual" Table

IP	Timestamp	Status
127.0.0.1	17/Oct/2025...	200

## Learn REGEX:

Basic Training: [https://www.w3schools.com/python/python\\_regex.asp](https://www.w3schools.com/python/python_regex.asp)

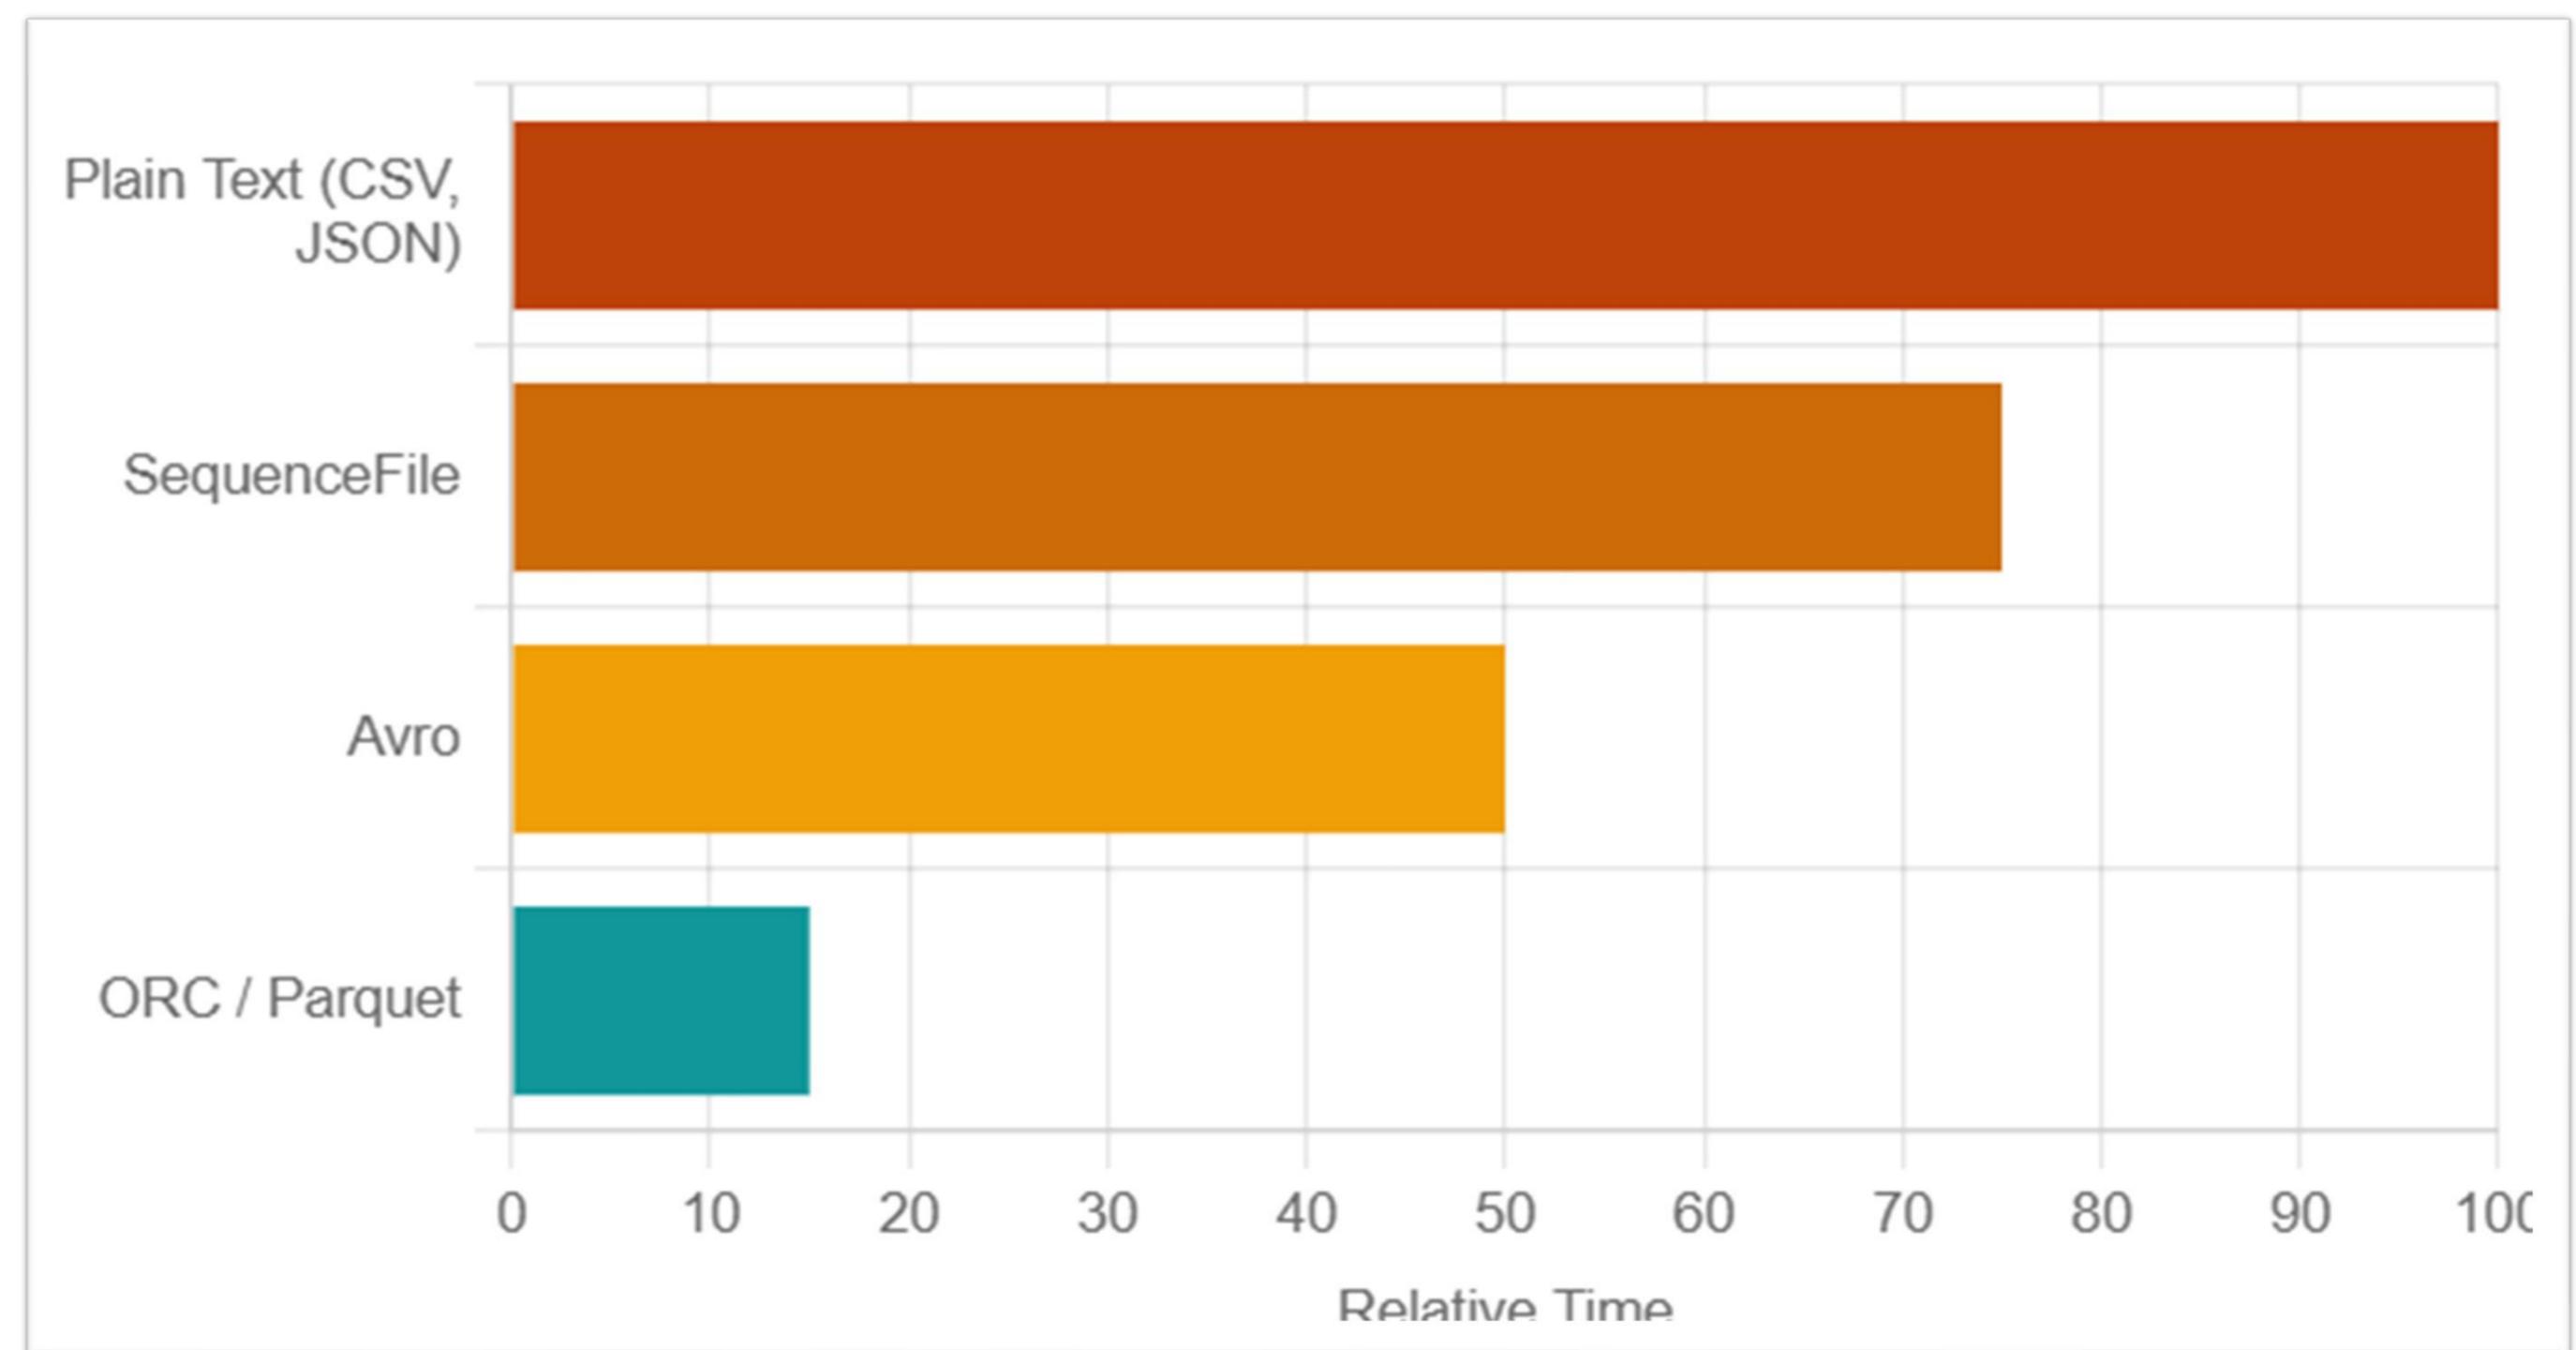
YouTube: <https://youtu.be/wx-SuoZXtuQ>



University of  
East London

# File Format Performance Impact

- After initial processing, converting to optimized columnar formats like Parquet or ORC dramatically speeds up subsequent analytical queries.
- Add **STORED AS PARQUET;** at the end of **CREATE TABLE intermediate\_logs\_parquet()**



# Advanced Hive Techniques & Tricks

Optimize productivity with powerful Hive features designed for complex data.

## Partitioning

Divides tables into smaller parts based on column values (e.g., date, region). This prunes data during queries, drastically reducing scan time and improving performance.

## Bucketing

Decomposes partitions into more manageable chunks based on a hash function of a column. Optimizes joins by pre-sorting and organizing data.

## Built-in Functions (UDFs)

Leverage a rich library of functions for string manipulation, date functions, and complex data type parsing (e.g., `get\_json\_object`, `xpath`) to extract features from unstructured data directly in your query.



University of  
East London

# Partitioning

Divides tables into smaller parts based on column values (e.g., date, region). This reduces scan time and improves query performance.

<b>user_id</b>	<b>name</b>	<b>region</b>
1	Alice	US
2	Bob	UK
3	Charlie	UK
4	Diana	UK
5	Eve	US

**US Partition**

User ID	Name	Region
1	Alice	US
5	Eve	US

```
-- Create a partitioned table
CREATE TABLE users_partitioned (
    user_id INT,
    name STRING
)
PARTITIONED BY (region STRING);

-- Insert data into the US partition
INSERT INTO TABLE users_partitioned
PARTITION (region='US')
VALUES
    (1, 'Alice'),
    (5, 'Eve');

-- Insert data into the UK partition
INSERT INTO TABLE users_partitioned
PARTITION (region='UK')
VALUES
    (2, 'Bob'),
    (3, 'Charlie'),
    (4, 'Diana');
```

**UK Partition**

User ID	Name	Region
2	Bob	UK
3	Charlie	UK
4	Diana	UK



# Bucketing

Bucketing is a method of dividing data into a fixed number of buckets (subsets) based on a hashing function or a specific range.

## Bucket 1

User ID	Name	Region
2	Bob	UK
4	Diana	UK

## Bucket 2

User ID	Name	Region
1	Alice	US
3	Charlie	UK
5	Eve	US

```
-- Create a bucketed table
CREATE TABLE users_bucketed (
    user_id INT,
    name STRING,
    region STRING
)
CLUSTERED BY (user_id)
INTO 2 BUCKETS;

-- Insert data into the bucketed table
INSERT INTO TABLE users_bucketed
VALUES
    (1, 'Alice', 'US'),
    (2, 'Bob', 'UK'),
    (3, 'Charlie', 'UK'),
    (4, 'Diana', 'UK'),
    (5, 'Eve', 'US');
```



# UDFs

Have a rich library of functions for string manipulation, date parsing, and unstructured data processing.

```
// Import the UDF class from Hive
import org.apache.hadoop.hive.ql.exec.UDF;

public class UpperCaseUDF extends UDF {
    public String evaluate(String input) {
        if (input == null) {
            return null;
        }
        return input.toUpperCase();
    }
}
```

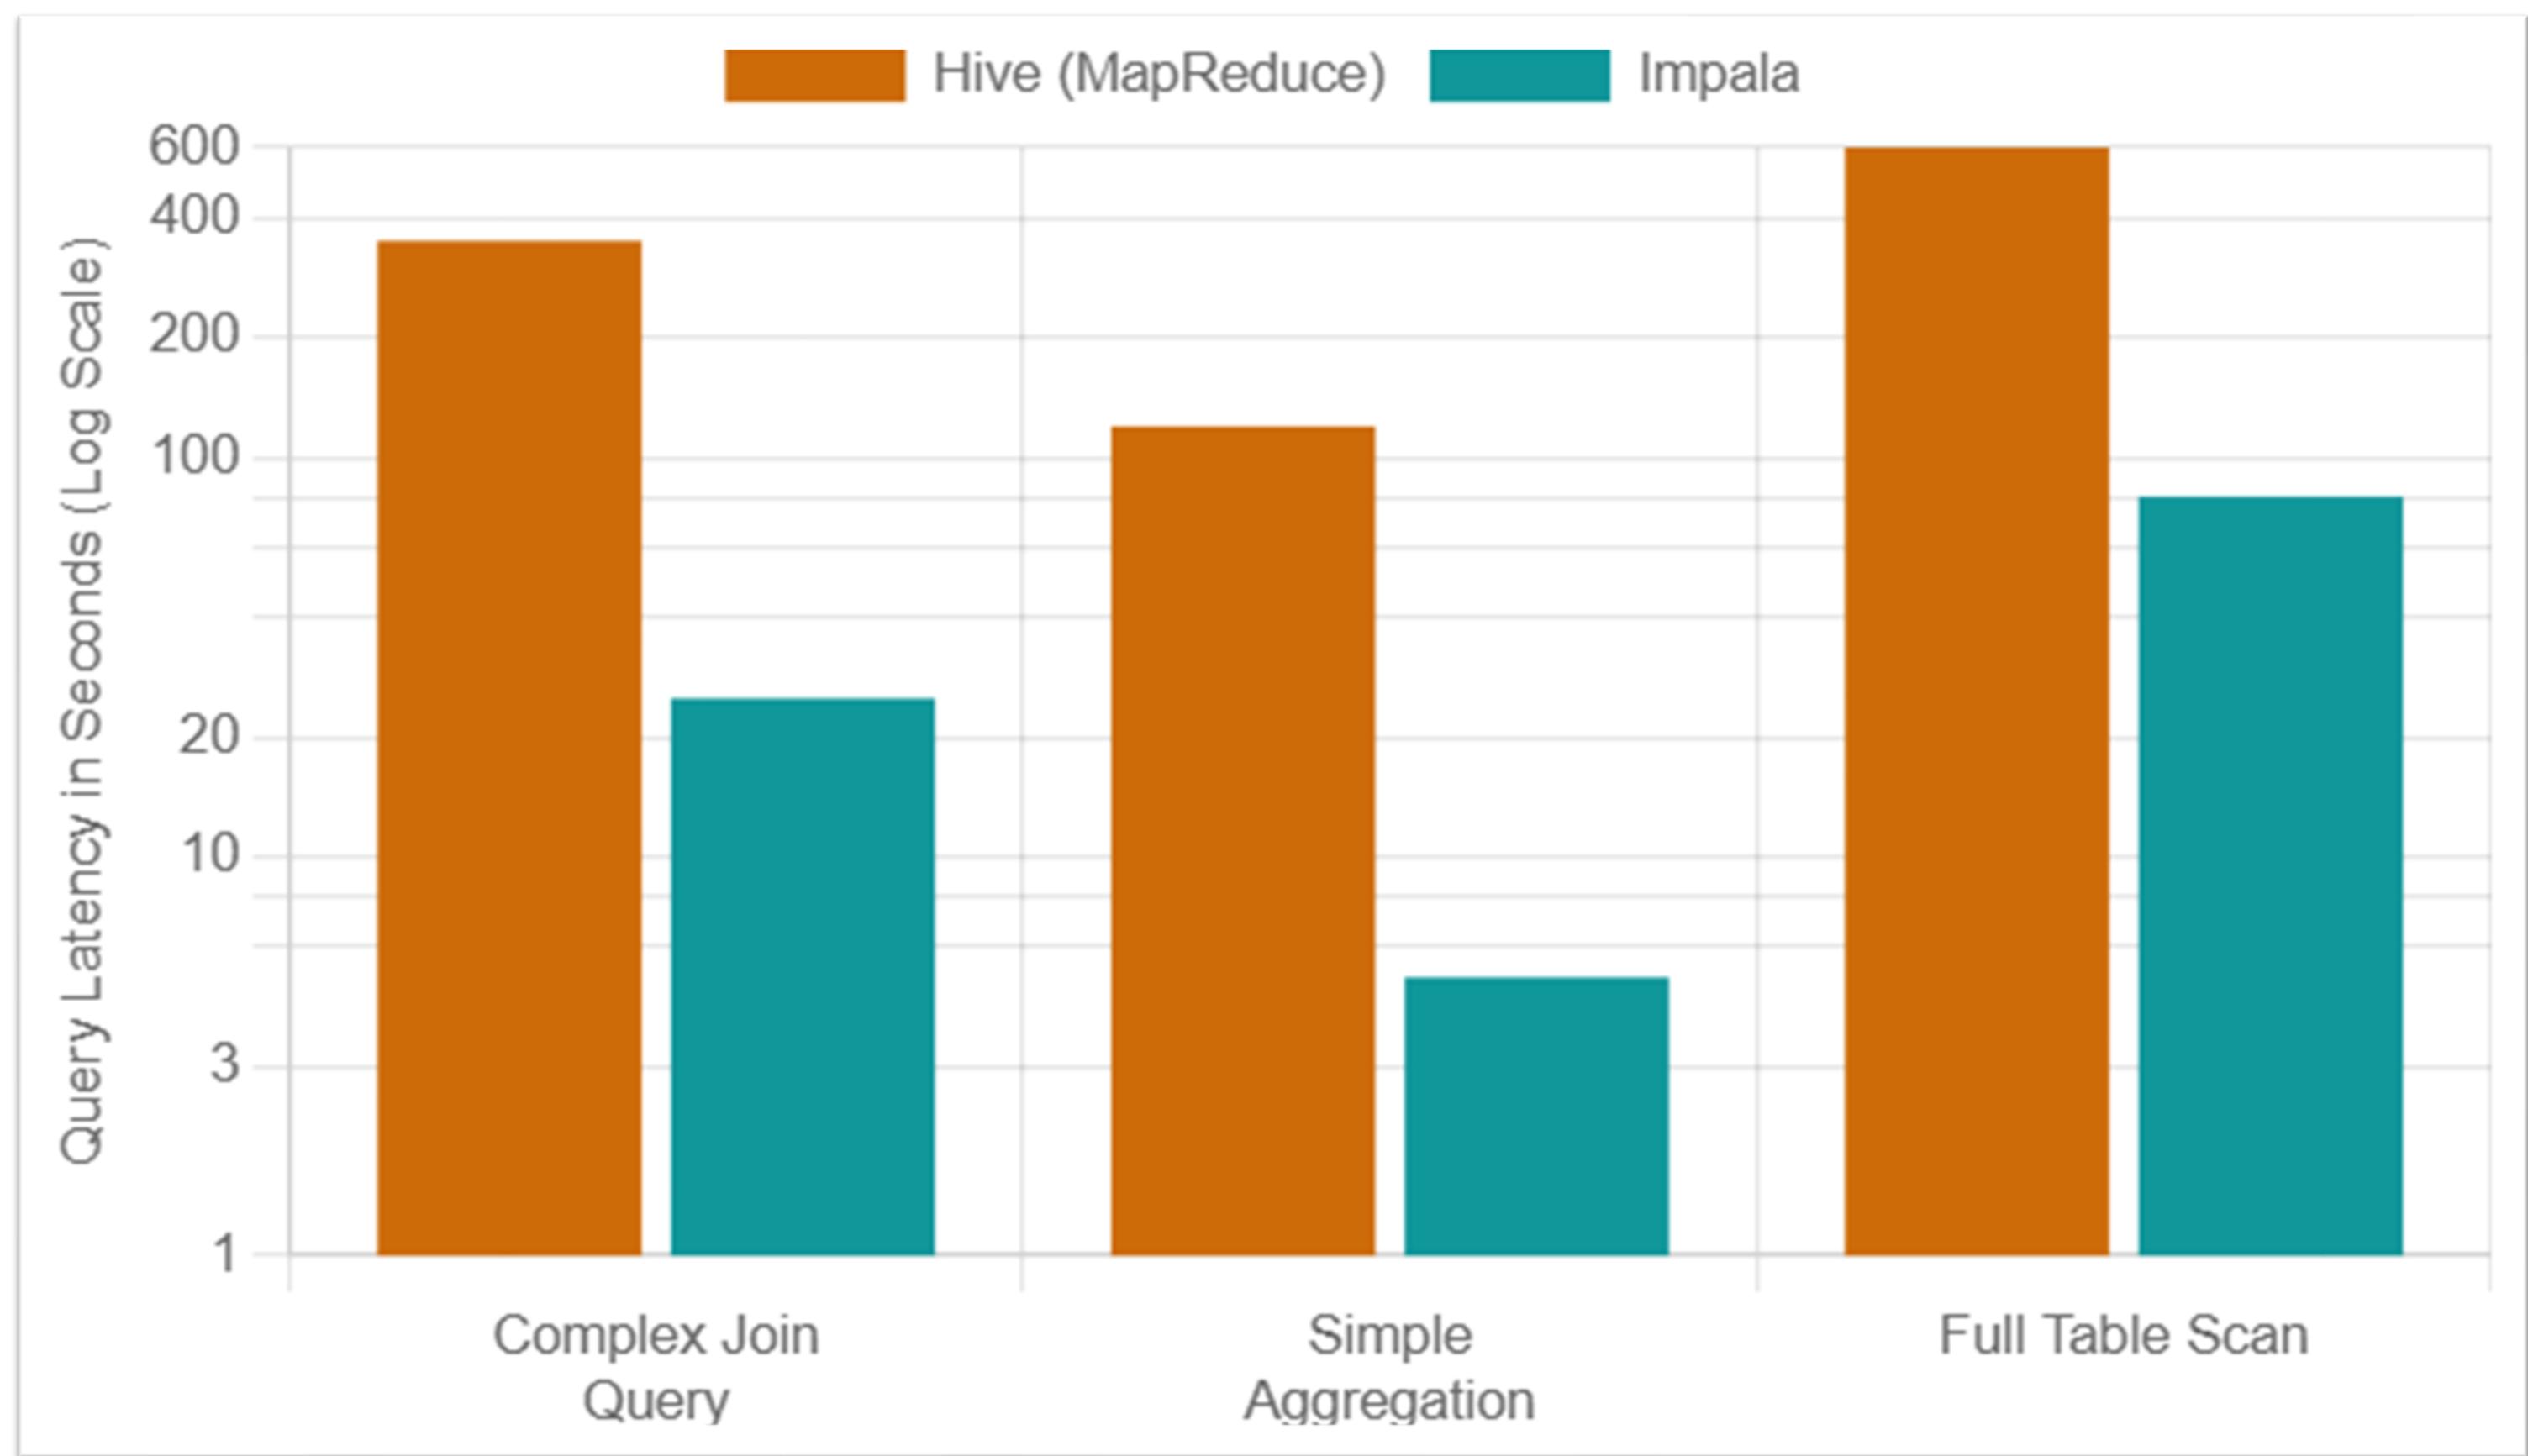
```
SELECT
    user_id,
    to_uppercase(user_name) AS uppercased_name
FROM user_profiles;
```



University of  
East London

# Accelerating Queries with Impala

Impala's in-memory, distributed query engine provides significantly lower latency compared to Hive's batch-oriented MapReduce approach.



# Tutorial 5

## Working with Unstructured Data in Hadoop

- Link to video: <https://youtu.be/ZM87yiOrUZ8>



University of  
East London

# Summary

- Discussed the challenges of Unstructured Data
- Discussed REGEX/Hive/Impala for Unstructured Data
- Discussed how to make queries fast
- Discussed Advanced Hive Optimizations (Partitioning, Bucketing, UDFs)